

Google Cloud Platform WordCount Search products and resources

Google account: Brendan Russ (bvr6@pitt.edu)

Home Marketplace Billing API APIs & Services Support IAM & Admin Getting started Security Anthos

COMPUTE App Engine Compute Engine

DASHBOARD ACTIVITY RECOMMENDATIONS

Project info

- Project name: WordCount
- Project ID: wordcount-293603
- Project number: 109974008917

[ADD PEOPLE TO THIS PROJECT](#)

[Go to project settings](#)

Resources

- Compute Engine: 3 instances
- Storage: 6 buckets

Compute Engine

CPU (%)

100% 80% 60% 40% 20% 0

10:30 11 PM

instance/cpu/utilization: 4.07%

[Go to Compute Engine](#)

Google Cloud Platform status

All services normal

[Go to Cloud status dashboard](#)

Billing

Estimated charges: USD \$0.00
For the billing period Oct 1 – 25, 2020

[View detailed charges](#)

Monitoring

Set up alerting policies

Create uptime checks

APIs

Requests (requests/sec)

```

bvr6@cluster-adv6-m:~/CloudWordCount$ hadoop jar WordCount.jar WordCount testData out put
Execution time: 1666milliseconds
20/10/26 03:22:09 INFO client.RMProxy: Connecting to ResourceManager at cluster-adv6-m/10.150.0.4:8032
20/10/26 03:22:09 INFO client.AHSProxy: Connecting to Application History server at cluster-adv6-m/10.150.0.4:10200
20/10/26 03:22:10 INFO input.FileInputFormat: Total input files to process : 1
20/10/26 03:22:10 INFO mapreduce.JobSubmitter: number of splits:1
20/10/26 03:22:10 INFO Configuration.deprecation: yarn.resourcemanager.system-metrics-publisher.enabled is deprecated. Instead, use yarn.system-metrics-publisher.enabled
20/10/26 03:22:10 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1603677175143_0002
20/10/26 03:22:11 INFO impl.YarnClientImpl: Submitted application application_1603677175143_0002
20/10/26 03:22:11 INFO mapreduce.Job: The url to track the job: http://cluster-adv6-m:8088/proxy/application_1603677175143_0002/
20/10/26 03:22:11 INFO mapreduce.Job: Running job: job_1603677175143_0002
20/10/26 03:22:18 INFO mapreduce.Job: Job job_1603677175143_0002 running in uber mode : false
20/10/26 03:22:18 INFO mapreduce.Job: map 0% reduce 0%
20/10/26 03:22:23 INFO mapreduce.Job: map 100% reduce 0%
20/10/26 03:22:31 INFO mapreduce.Job: map 100% reduce 20%
20/10/26 03:22:32 INFO mapreduce.Job: map 100% reduce 40%
20/10/26 03:22:33 INFO mapreduce.Job: map 100% reduce 60%
20/10/26 03:22:35 INFO mapreduce.Job: map 100% reduce 100%
20/10/26 03:22:36 INFO mapreduce.Job: Job job_1603677175143_0002 completed successfully
20/10/26 03:22:36 INFO mapreduce.Job: Counters: 51
  File System Counters
    FILE: Number of bytes read=2082
    FILE: Number of bytes written=1250533
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=1558
    HDFS: Number of bytes written=1393
    HDFS: Number of read operations=28
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=15
  Job Counters
    Killed reduce tasks=1
    Launched map tasks=1
    Launched reduce tasks=5
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=13328
    Total time spent by all reduces in occupied slots (ms)=134736
    Total time spent by all map tasks (ms)=3332
    Total time spent by all reduce tasks (ms)=33684
    Total vcore-milliseconds taken by all map tasks=3332

```

```
Total time spent by all reduces in occupied slots (ms)=134736
Total time spent by all map tasks (ms)=3332
Total time spent by all reduce tasks (ms)=33684
Total vcore-milliseconds taken by all map tasks=3332
Total vcore-milliseconds taken by all reduce tasks=33684
Total megabyte-milliseconds taken by all map tasks=6823936
Total megabyte-milliseconds taken by all reduce tasks=68984832
Map-Reduce Framework
  Map input records=23
  Map output records=253
  Map output bytes=2446
  Map output materialized bytes=2082
  Input split bytes=114
  Combine input records=253
  Combine output records=165
  Reduce input groups=165
  Reduce shuffle bytes=2082
  Reduce input records=165
  Reduce output records=165
  Spilled Records=330
  Shuffled Maps =5
  Failed Shuffles=0
  Merged Map outputs=5
  GC time elapsed (ms)=731
  CPU time spent (ms)=4220
  Physical memory (bytes) snapshot=1563582464
  Virtual memory (bytes) snapshot=20804145152
  Total committed heap usage (bytes)=1240989696
Shuffle Errors
  BAD ID=0
  CONNECTION=0
  IO ERROR=0
  WRONG LENGTH=0
  WRONG MAP=0
  WRONG REDUCE=0
WordCount$TokenizerMapper$CountersEnum
  INPUT WORDS=253
File Input Format Counters
  Bytes Read=1444
File Output Format Counters
  Bytes Written=1393
```

10/11/16
10:11:16
10:11:16